



David Mohaisen, Professor

University of Central Florida

Department of Computer Science

Cybersecurity and Privacy Cluster

<https://cs.ucf.edu/~mohaisen/>

Advanced Word Embeddings and Evaluation

- ① Gradient descent, revisited
- ② Skip-gram with negative sample
- ③ Window-based representation
- ④ Count-based vs. direct prediction
- ⑤ Embedding meanings, word sense
- ⑥ How to evaluate word vectors

Gradient Descent

- $f(x) = x^2, f'(x) = 2x, f'(0) = 0.$
- Initial guess $x = 4; f'(x) = 2*4 = 8$
- Update: $x_{new} = x_{old} - \alpha f'(x_{old})$
 $x_{new} = 4 - 0.1 * 8 = 3.2$
 $f(3.2) = 2*3.2 = 6.4$
- ...
- Repeat: $x_{new} = 3.2 - 0.1*6.4 = 2.56; f(x) = 2*2.56=5.12$
- ...

- Update equation (in matrix notation):

$$\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$$

$\alpha = \text{step size or learning rate}$

- Update equation (for single parameter):

$$\theta_j^{new} = \theta_j^{old} - \alpha \frac{\partial}{\partial \theta_j^{old}} J(\theta)$$

Gradient Descent

- Define variables:
 - Predictions (vector of predicted values):
 - Target (actual values):
 - The function is MSE loss; $n = 3$:

$$\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \hat{y}_3]$$

- Compute derivative w.r.t. $\mathbf{y} = [y_1, y_2, y_3]$.

$$L(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

- Substitute values; let's say:

- Compute difference:
- Compute the gradient:

$$\frac{\partial L}{\partial \hat{\mathbf{y}}} = \frac{2}{n} (\hat{\mathbf{y}} - \mathbf{y})$$

$$\hat{\mathbf{y}} = [2.0, 3.0, 4.0]$$

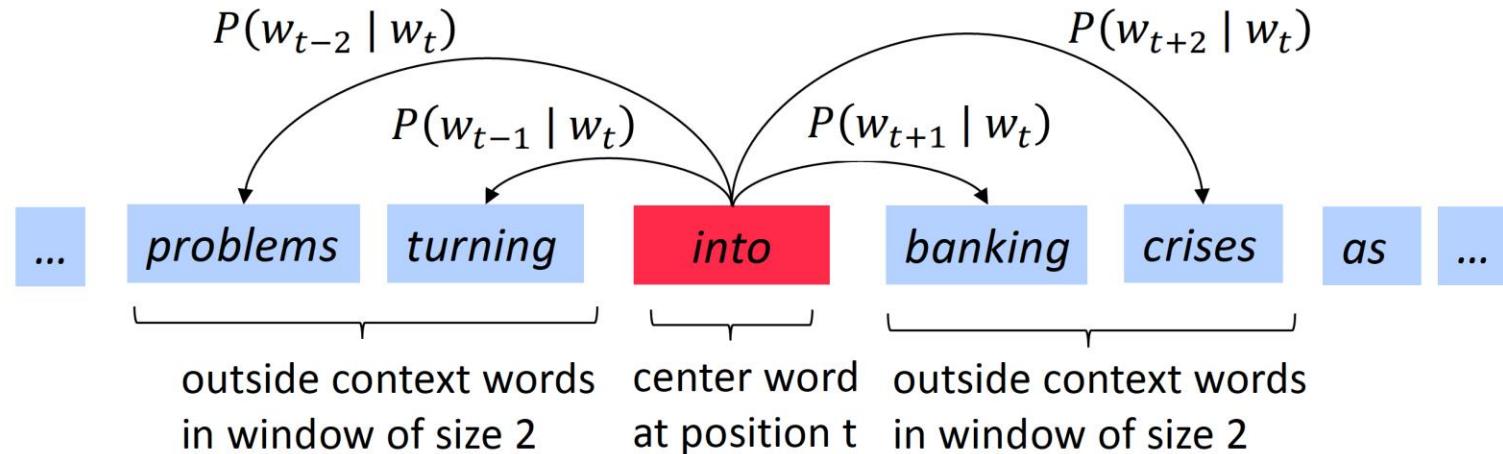
$$\mathbf{y} = [1.5, 3.5, 4.5]$$

$$\hat{\mathbf{y}} - \mathbf{y} = [2.0 - 1.5, 3.0 - 3.5, 4.0 - 4.5] = [0.5, -0.5, -0.5]$$

$$\frac{\partial L}{\partial \hat{\mathbf{y}}} = \frac{2}{3} [0.5, -0.5, -0.5] = \left[\frac{1}{3}, -\frac{1}{3}, -\frac{1}{3} \right]$$

Summary of Word2Vec

- Iterate over each word for the whole corpus
- Predict surrounding words using word vectors



$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

- Update vectors so you can predict well

Stochastic Gradient Descent

- Problem: $J(\boldsymbol{\theta})$ is a function of all windows in the corpus (potentially billions)
 - So the update of $J(\boldsymbol{\theta})$ over $\boldsymbol{\theta}$ is expensive
- You would wait a very long time before making a single update.
 - Very bad idea for pretty much all problems
- Solution: SGD: repeatedly sample windows, and update after each one.

Skip-gram with Negative Sample

- The normalization is too expensive

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

- Hence, we implement the skip-gram
- Idea: train a binary logistic regression for a true pair (center with actuals) versus several noise pairs (center with random words)

Skip-Gram with Negative Sampling

- Overall objective function to maximize

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T J_t(\theta)$$

$$J_t(\theta) = \log \sigma(u_o^T v_c) + \sum_{i=1}^k \mathbb{E}_{j \sim P(w)} [\log \sigma(-u_j^T v_c)]$$

- Meaning: we maximize two words co-occurring in first log.

Skip-Gram: Procedure

- We take k negative samples (using the word probabilities)
- Maximize the probability that real outside word appears, minimize prob. that random words appear around the center word.
- Why not capture co-occurrence directly:
 - Window: similar to word2vec, use window around each word \Leftrightarrow capture both syntactic and semantic information.
 - Word-document co-occurrence matrix gives general topics (all sports terms will have similar entities) leading to "latent semantic analysis".

Skip-gram with Negative Sample: Details

- **Objective:** Learn dense word embeddings by predicting context words given a target word.
- **How It Works:**
 1. **Input:** Target word and its context window.
 2. **Positive Samples:** Word pairs from the corpus ($w_{target}, w_{context}$).
 3. **Negative Samples:** Random pairs (w_{target}, w_{noise}) generated from a noise distribution.
- 1. **Training:** Binary classification to distinguish real (positive) from random (negative) pairs.
- 2. **Output:** Embeddings that capture semantic similarity.
- **Advantages:**
 - Efficient for large vocabularies.
 - Captures semantic relationships (e.g., analogies, similarities).
 - Widely used in NLP tasks (e.g., translation, sentiment analysis).
- **Applications:** Pre-trained embeddings for downstream NLP tasks, semantic search, etc.

Example: Window-Based

- Window length = 1 (common 5-10)
- Symmetric (irrelevant whether left or right context).
- Example corpus:
 - I like deep learning
 - I like NLP
 - I enjoy flying

Example: Window-Based, cont.

- Example corpus:
 - I like deep learning
 - I like NLP
 - I enjoy flying

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Window-Based, cont.

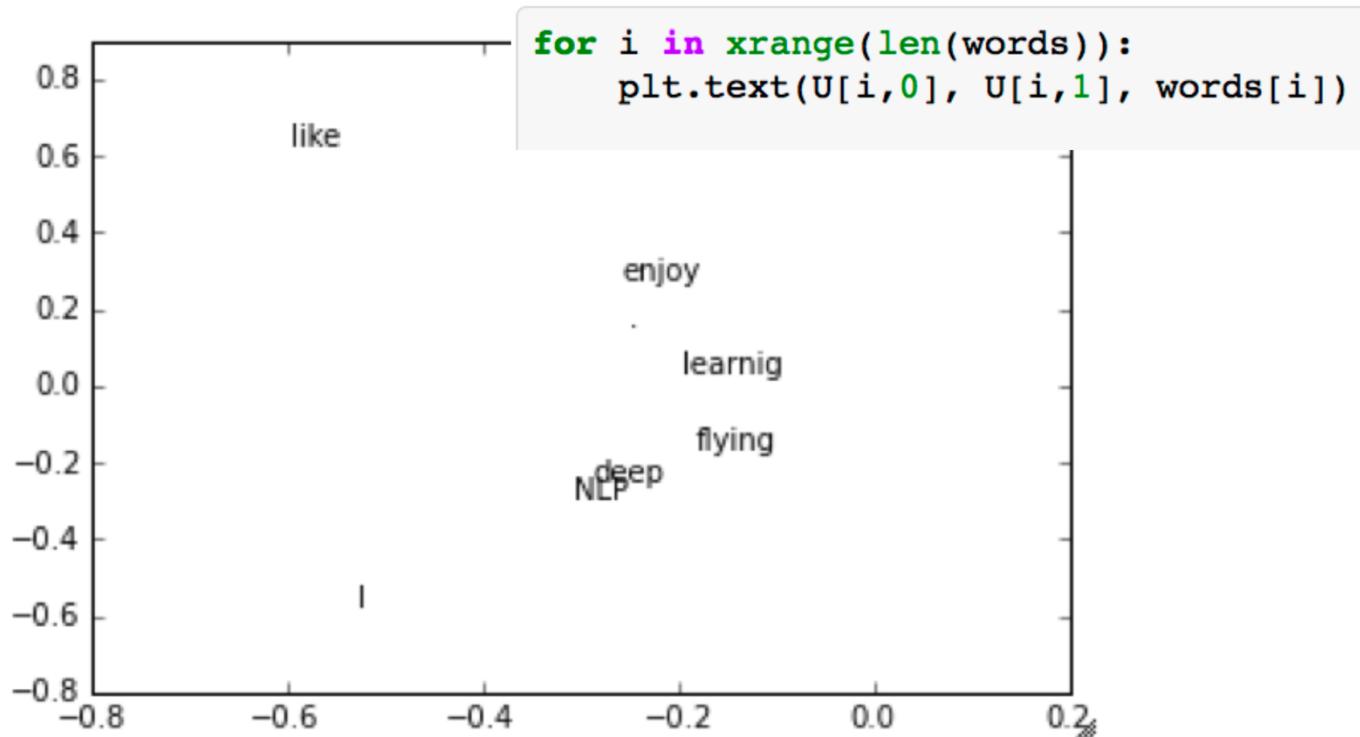
- Solution: low dimensional vectors
 - Store most of the important info in a fixed small number of dimensions: dense vector
 - Usually 25-1000 dimensions, similar to word2vec
 - How to reduce the dimensionality?
 - Singular value decomposition (SVD) of co-occurrence matrix X ; factor X in to USV^T such that U and V are orthogonal. Retain k singular values to generalize.
 - Expensive to compute for large matrices.

Example in Python

```
import numpy as np
la = np.linalg
words = ["I", "like", "enjoy",
          "deep", "learnig", "NLP", "flying", ".."]
X = np.array([[0,2,1,0,0,0,0,0],
              [2,0,0,1,0,1,0,0],
              [1,0,0,0,0,0,1,0],
              [0,1,0,0,1,0,0,0],
              [0,0,0,1,0,0,0,1],
              [0,1,0,0,0,0,0,1],
              [0,0,1,0,0,0,0,1],
              [0,0,0,0,1,1,1,0]])
U, s, Vh = la.svd(X, full_matrices=False)
```

Example: Visualization

- Printing first two columns of U corresponding to the biggest singular values



Hacks to X

- Scaling the count helps a lot:
 - Problem: function words (the, he, has, a, etc.) are too frequent \Leftrightarrow syntax has too much impact. Some fixes:
 - $\min(X, t)$, with $t=100$
 - Ignore them all
 - Ramped windows that count closer words more
 - Use Pearson correlations instead of counts, then set negative values to 0
 - Other approaches?

Count-Based vs. Direct Prediction

Count-based

- Fast training
- Efficient use of stats
- Primarily used to capture word similarity
- Disproportionate importance given to large counts

Direct-Prediction

- Scales with corpus size
- Inefficient usage of statistics
- Generate improved performance on other tasks
- Can capture complex patterns beyond word similarity

Encoding Meaning

- Ratios of co-occurrence probabilities can encode meaning components

	$x = \text{solid}$	$x = \text{gas}$	$x = \text{water}$	$x = \text{random}$
$P(x \text{ice})$	large	small	large	small
$P(x \text{steam})$	small	large	large	small
$\frac{P(x \text{ice})}{P(x \text{steam})}$	large	small	~1	~1

Encoding Meanings

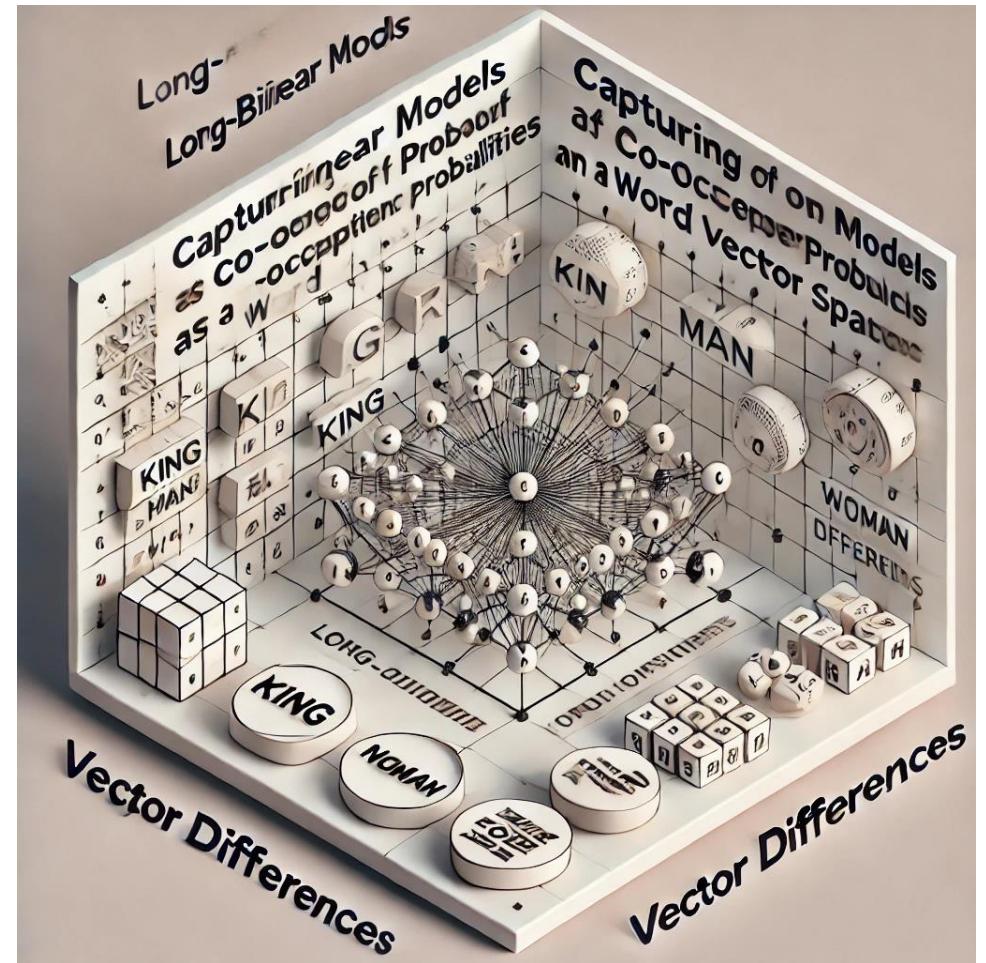
- How can we capture ratios of co-occurrence probabilities as a linear meaning components in a word vector space?

A: Log-bilinear model:

$$w_i \cdot w_j = \log P(i|j)$$

with vector differences

$$w_x \cdot (w_a - w_b) = \log \frac{P(x|a)}{P(x|b)}$$



How to Evaluate Word Vectors?

- Intrinsic:
 - Specific subtask
 - Fast to compute
 - Help understand the system
 - Not clear if real helpful unless correlation to real task is established
- Extrinsic
 - Evaluation on a real task
 - Can take long time to compute
 - Unclear if subsystem is the problem or its iterations
 - If replacing exactly one subsystem with another improves accuracy \Leftrightarrow winning

Intrinsic Evaluation

- **Focus:** Measures the quality of embeddings on tasks that directly test word similarity or relatedness.
- **Methods:**
 - **Word Similarity Tasks:** (1) Use human-annotated datasets like WordSim-353 or SimLex-999. (2) Evaluate cosine similarity b/w vectors and compare against human judgments.
 - **Analogy Tasks:** (2) Solve analogies like "king - man + woman = queen." (2) Popular dataset: Google Analogy Dataset.
 - **Clustering:** Group similar words together and measure coherence.
 - **Qualitative Analysis:** Visualize embeddings using techniques like t-SNE or PCA.
- **Advantages:** (1) Quick and interpretable. (2) Independent of specific NLP applications.
- **Disadvantages:** May not reflect performance on downstream tasks.

Extrinsic Evaluation

- **Focus:** Measures the performance of word embeddings in real-world NLP applications.
 - **Methods:**
 - **Text Classification:** Use embeddings as input features for tasks like sentiment analysis or spam detection.
 - **Named Entity Recognition (NER):** Assess embeddings in identifying entities (e.g., names, dates).
 - **Machine Translation:** Evaluate their role in generating accurate translations.
 - **Question Answering:** Test embeddings in retrieving correct A to Q.
 - **Adv.:** Measures practical utility in real-world tasks.
 - **Disadvantages:** Computationally expensive and task-specific.



Word Sense

- **Words Have Multiple Meanings**
 - Ambiguity is common, especially for frequent or historical words.
 - Example: "Pike" (fish, weapon, road).
- **Challenge with Word Embeddings**
 - Static embeddings (e.g., Word2Vec) assign one vector per word.
 - Do they capture all meanings, or create a mess?
- **Solution: Contextual Embeddings**
 - Models like BERT assign different vectors based on context.
 - Helps resolve ambiguity dynamically.

Word Sense, example

- Word: "Bank" – 5 Different Meanings

1. Financial Institution

1. *"I deposited money in the bank."*
2. Meaning: A place to store and manage money.

2. Riverbank

1. *"He sat on the bank of the river."*
2. Meaning: The land beside a river.

3. Action of Banking (Tilting)

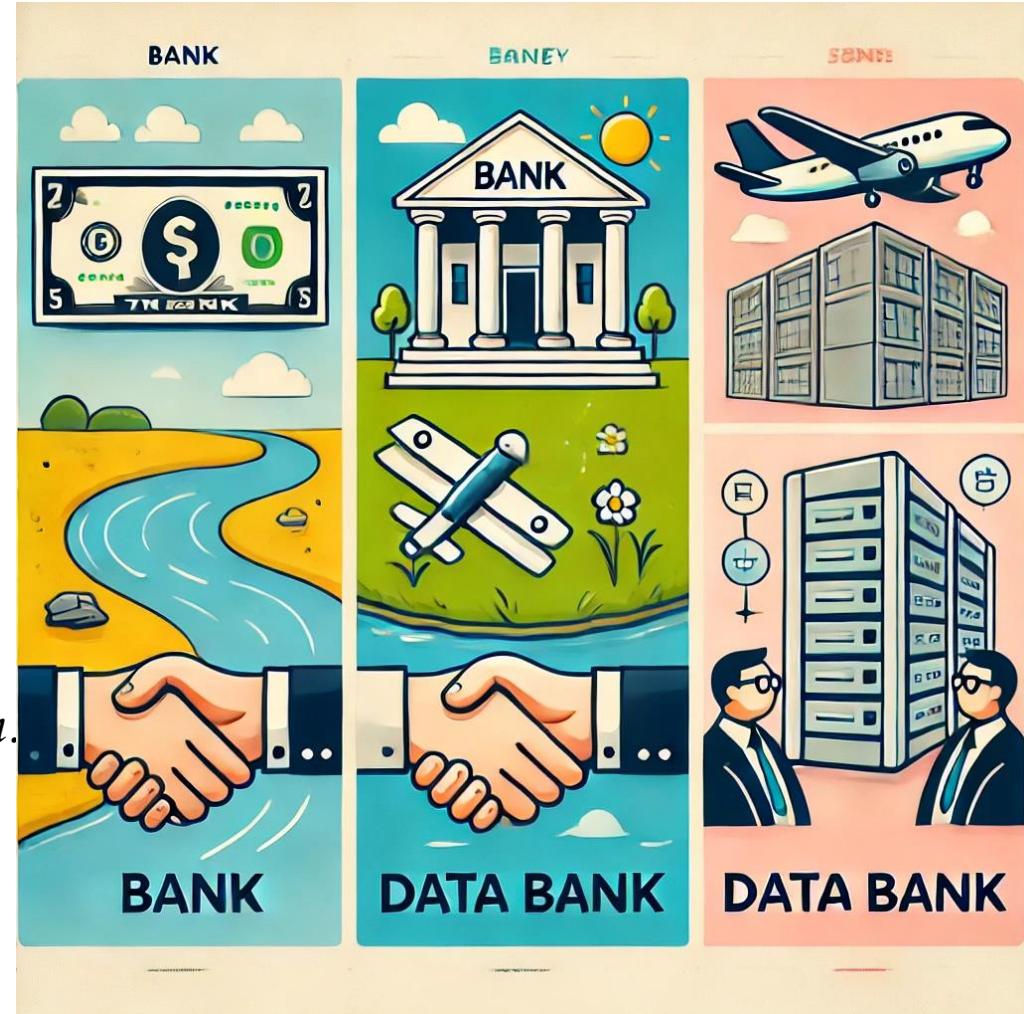
1. *"The airplane banked sharply to the left."*
2. Meaning: To tilt or turn while in motion.

4. Storage (e.g., Data Bank)

1. *"The company maintains a bank of customer information."*
2. Meaning: A collection or repository.

5. Verb: Trust or Rely

1. *"You can bank on her to deliver the project."*
2. Meaning: To trust or depend on someone.



Word Sense, Another Example (Pike)

- A sharp point or staff
- A type of elongated fish
- A railroad line or system
- A type of road
- The future (coming down the pike)
- A type of body position (as in diving)
- To kill or pierce with a pike
- To make one's way (pike along)
- In Australian English, pike means to pull out from doing something: I reckon he could have climbed that cliff, but he piked!

Global Context Solutions

- **Attention Mechanisms:**
 - Focus on relevant words across the text.
- **Contextual Embeddings:**
 - Models like BERT/GPT use full text for word meaning.
- **Topic Models:**
 - Capture document themes for better representation.



Improving Word with Global Context

- Idea: cluster word windows around words, retain with each word assigned multiple different clusters.

