David Mohaisen, Professor
University of Central Florida
Department of Computer Science
Cybersecurity and Privacy Cluster

https://cs.ucf.edu/~mohaisen/

Reasoning and Generalization in Language Models

① Cognitive differences
② Reasoning
③ Language models simulate reasoning
④ CoT
⑤ Open challenges
⑥ Summary and Takeaways

Objectives

- Understand the **cognitive differences** between humans and machines in learning and reasoning.
- Explore the concept of **reasoning** and its forms: deductive, inductive, abductive.
- Examine how **language models simulate reasoning**, including strengths and limitations.
- Analyze techniques like **Chain-of-Thought (CoT) prompting** and their impact on LLM performance.
- Evaluate evidence for reasoning in LLMs through benchmarks and behavioral studies.
- Identify **open challenges and future directions** in building reasoning-capable AI systems.

Humans vs. Machines

Humans vs. Machines — Cognitive Abilities

Objective: Understand the fundamental differences in learning, generalization, and reasoning between humans and modern machine learning systems.

- **Human Intelligence**:
  - Adaptable to novel tasks with minimal supervision.
  - Capable of abstract reasoning and contextual understanding.
  - Intrinsically interpretable decision-making.
- **Machine Intelligence**:
  - Relies on large-scale data and supervised training.
  - Susceptible to overfitting and domain-specificity.
  - Lacks innate transparency and cognitive reasoning.
- **Key Question**: *Why do machines struggle with generalization and reasoning despite surpassing humans in raw computational power?*

Few-Shot Learning, Minimal Examples

Human Capability:

- Can learn complex concepts from a handful of examples.
- Leverages prior knowledge and analogical reasoning.
- Incorporates multimodal cues and semantic priors.

Machine Limitation:

- Few-shot learning remains an open challenge in ML.
- Algorithms like meta-learning and prompt-based LLMs attempt to emulate this.
- Performance often degrades in low-data or unseen classes.

Example:

- Human: Learns the concept of a "giraffe on a surfboard" from one image.
- Machine: Requires extensive data augmentation and fine-tuning.

OOD Robustness, Dealing with the Unknown

Human Capability:

- Naturally robust to distribution shifts (e.g., visual noise, contextual ambiguity).
- Can infer and act reasonably in unfamiliar or adversarial scenarios.

Machine Limitation:

- Deep neural networks are sensitive to shifts in data distribution.
- Failure in OOD settings exposes vulnerability to adversarial attacks and poor generalization.

Technical Insight:

- Distributional shift violates i.i.d. assumptions foundational to supervised learning.
- Techniques such as domain generalization, invariant representation learning, and test-time adaptation are ongoing research areas.

Explainability, From Intuition to Attribution

Human Capability:

- Can articulate reasons behind a decision or prediction.
- Integrates introspection, context, and commonsense reasoning.

Machine Limitation:

- Neural models are often viewed as “black boxes”.
- Post-hoc methods (e.g., SHAP, LIME) provide limited and sometimes inconsistent explanations.
- Faithfulness and causality in explanations remain unsolved.

Emerging Directions:

- Symbolic reasoning overlays.
- Neuro-symbolic models.
- Causal interpretability frameworks.

Reasoning, Missing Ingredient in Machine Intel.

Human Capability:

- Engages in multi-step, causal, and counterfactual reasoning.
- Flexible across deductive, inductive, and abductive paradigms.

Machine Limitation:

- Most models rely on statistical pattern matching, not logical inference.
- Reasoning often requires explicit scaffolding or structured prompts.

Notable Progress:

- Chain-of-Thought prompting in LLMs.
- Neuro-symbolic architectures integrating logic and learning.
- Differentiable reasoning modules.

Comparison of Humans and Machines

Capability	Humans	Machines
Few-Shot Learning	Strong ability to learn from few examples	Typically requires large datasets and fine-tuning
Out-of-Distribution Robustness	Generalizes well to novel or unseen scenarios	Often fails or degrades under distributional shifts
Explainability	Can articulate and justify decisions	Often operates as a "black box" with limited interpretability
Reasoning	Engages in abstract, causal, and logical reasoning	Lacks native reasoning; depends on engineered mechanisms
Data Efficiency	Learns efficiently using prior knowledge and context	Data-hungry and brittle without large-scale training
Adaptability	Flexible across tasks and modalities	Typically narrow and task-specific
Contextual Understanding	Rich context integration via multimodal and experiential input	Limited and superficial contextual grasp

Summary and Discussion

Key Takeaways:

- Humans excel in few-shot generalization, reasoning, and explainability.
- Machine learning models lag in adaptability, robustness, and cognitive transparency.
- Bridging the gap requires rethinking data-centric paradigms, integrating reasoning structures, and prioritizing interpretability.

Discussion:

- Can AI systems become better at reasoning without symbolic representations?
- What benchmarks can fairly compare human and machine intelligence?
- How do large language models change the landscape of human-machine comparisons?

Reasoning

What is Reasoning?

Definition: Reasoning is the **cognitive process of drawing conclusions**, making decisions, or solving problems through logical thinking and evidence-based analysis.

Key Characteristics:

- Involves **structured thought**, often beyond immediate sensory input.
- Employs **evidence, prior knowledge, and logical principles**.
- Bridges the gap between perception and decision-making.

Type	Description	Example
Deductive	Deriving specific conclusions from general premises.	*All humans are mortal → Socrates is mortal.*
Inductive	Generalizing from specific observations to broader generalizations.	*Observed 100 swans → All swans are white.*
Abductive	Inferring the most likely explanation from incomplete information.	*The grass is wet → It probably rained.*

Formal vs Informal Reasoning

Aspect	Formal Reasoning	Informal Reasoning
Structure	Rule-based, follows logical systems	Heuristic, often based on intuition
Context	Used in mathematics, logic, formal proofs	Used in everyday decision-making
Tools	Symbolic logic, theorem provers, axiomatic systems	Experience, analogies, mental models

Importance in AI:

- Core to building **interpretable, trustworthy, and generalizable** systems.
- Traditional AI used symbolic reasoning (e.g., expert systems).
- Modern ML often lacks explicit reasoning — leading to research on:
  - Chain-of-thought prompting
  - Neuro-symbolic models
  - Differentiable reasoning modules

Reasoning in Language Models

- Reasoning in language models (LMs) remains a **loosely defined** and often ambiguous concept.
- It typically refers to the model's ability to **draw conclusions or make inferences** based on textual input, resembling human reasoning processes.
- **Lack of a Formal Definition:**
  - No universally accepted definition of "reasoning" in the context of LMs.
  - Often conflated with **pattern recognition**, **inference**, or **structured output generation**.
  - Efforts to define reasoning tend to rely on the **model's output behavior** rather than internal mechanisms.

Informal Deductive Reasoning

A prevalent mode in LLM benchmarks and evaluations

Definition:

- Informal deductive reasoning refers to the process where a **conclusion logically follows from premises**, assuming the premises are true.
- Unlike formal deduction, it does not require symbolic representation or logical proofs but adheres to the **structure of valid reasoning**.

Premise 1: All mammals have lungs.
Premise 2: A dolphin is a mammal.
Conclusion: Therefore, a dolphin has lungs.

Challenges in Modeling Reasoning:

- **Surface-level correlations** vs. true logical reasoning.
- Prone to **hallucination**, contradiction, or invalid inference.
- Lack of **explicit control over the reasoning steps** or formal guarantees.

Emerging Techniques

Technique	Goal
Chain-of-Thought Prompting	Induce multi-step reasoning via examples
Self-consistency Decoding	Improve reasoning accuracy via sampling
Tool-Augmented Reasoning	Leverage calculators, search, or code
Formal Benchmarks (e.g., LogiQA, ProofWriter)	Test logical inference and symbolic deduction

CoT

Chain-of-Thought (CoT) Prompting

Definition:

- **Chain-of-Thought (CoT) Prompting** is a prompting technique that encourages a language model to **generate intermediate reasoning steps** before arriving at a final answer.
- The goal is to simulate **multi-step, interpretable reasoning** to improve performance on tasks requiring logic, arithmetic, or deduction.

Core Idea:

- Instead of directly predicting the answer, the model is prompted with examples that include **step-by-step reasoning paths**.
- This scaffolding allows the model to:
  - Break down complex problems,
  - Maintain context across steps,
  - Reduce reasoning errors.

Chain-of-Thought (CoT) Prompting

- **Example** (Arithmetic Reasoning)
  - **Model Output (with CoT prompting)**:

```
Q: If there are 3 cars and each car has 4 wheels, how many wheels are there in total?
A: Each car has 4 wheels. So, 3 cars have 3 × 4 = 12 wheels. Answer: 12
```

  - **Benefits**

```
Each car has 4 wheels. 3 cars means 3 × 4 = 12. Final answer: 12.
```

Aspect	Impact
Accuracy	Improves performance on multi-step reasoning tasks
Interpretability	Makes the reasoning process more transparent and debuggable
Generalization	Enhances transfer to unseen, logic-intensive tasks

Rationale Engineering

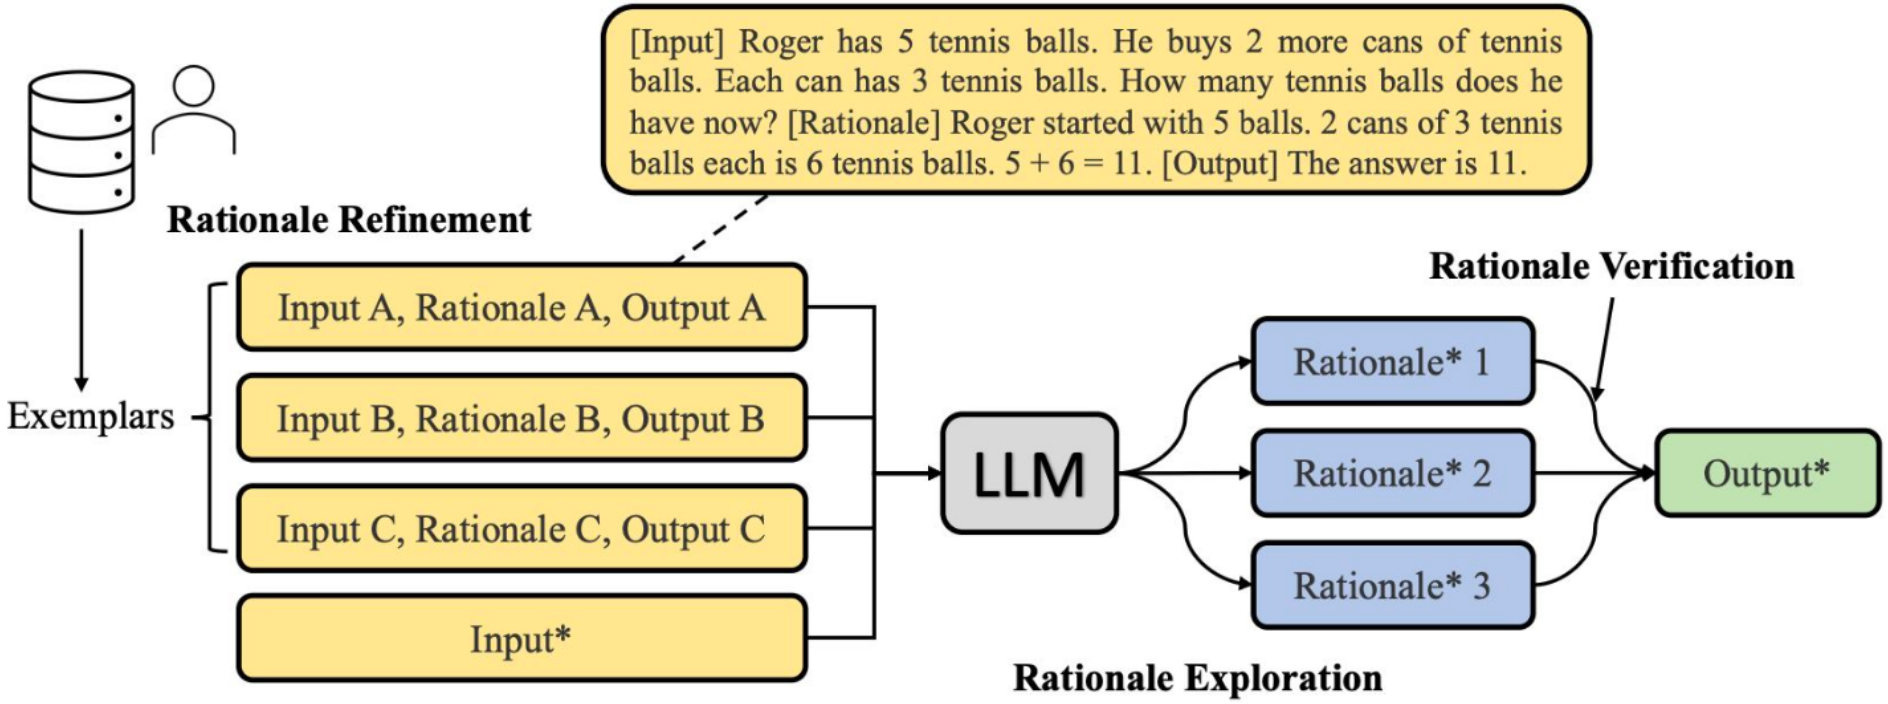

Jie Huang and Kevin Chen-Chuan Chang. Towards reasoning in large language models: A survey. 2023.

Challenges

Limitations and Variants

Limitations

- **Scaling requirement**: Effective primarily in large models (e.g., GPT-3 175B+).
- **Prompt sensitivity**: Quality of reasoning depends heavily on example phrasing.
- **Lack of guarantees**: Still susceptible to hallucination or faulty logic.

Variants and Enhancements:

Technique	Description
Self-Consistency	Sample multiple reasoning paths and aggregate final answers
Program-Aided CoT	Combine CoT with executable code snippets (e.g., Python)
CoT with Tool Use	Integrate external tools (calculators, databases)
Automatic CoT	Induce reasoning chains without manual exemplars

Are Language Models Really Able to Reason?

- **Central Question**: Can language models genuinely **perform reasoning**, or are they merely mimicking patterns observed in training data?
- Empirical Evidence Suggesting Reasoning Abilities in LLMs

Observation	Interpretation
High performance on reasoning benchmarks	LLMs achieve strong results on tasks such as GSM8K, ARC, and DROP
(Suzgun et al., 2022)	Indicates potential for multi-step reasoning under certain prompts
Chain-of-Thought (CoT) prompting improves reasoning	Step-by-step prompting reveals latent structured reasoning ability
(Wei et al., 2022)	Suggests models can learn to simulate logical chains of thought
Human-like biases in reasoning patterns	LLMs show content effects and belief bias similar to human cognition
(Dasgupta et al., 2022)	Implies overlap between human and LLM reasoning heuristics

Are Language Models Really Able to Reason?

- **Interpretive Perspectives**

Perspective	Claim
Skeptical View	LLMs mimic reasoning via pattern completion; no true understanding
Pragmatic View	If outputs are useful and interpretable, internal mechanisms are secondary
Cognitive Modeling View	LLMs replicate certain human reasoning biases and heuristics

- **Conclusion & Open Questions**
- There is growing evidence that **LLMs can perform forms of informal reasoning**.
- However, whether this counts as **"real reasoning"** depends on how reasoning is defined:
  - Symbolic inference?
  - Structured decision-making?
  - Human-like heuristics?
- Ongoing research seeks to disentangle **reasoning competence** from **statistical fluency**.

Model	Year	Parameters	Key Reasoning Technique	Strengths	Limitations
GPT-3	2020	175B	Few-shot prompting, basic CoT	General-purpose reasoning; strong few-shot performance	No explicit reasoning module; brittle to prompt variation
PaLM	2022	540B	Chain-of-Thought prompting	High performance on logic/math benchmarks; interpretable CoT	Requires manual CoT examples; needs very large model size
Minerva	2022	62B–540B	CoT + Scientific/Math fine-tuning	Excels at math and symbolic reasoning	Domain-specific; needs specialized data and tuning
Chinchilla	2022	70B	Data-efficient training, CoT capable	Outperforms larger models due to better compute/data balance	Less open usage; reasoning is implicit, not modular
ChatGPT (GPT-3.5)	2022	~175B (est.)	Instruction tuning + dialogue CoT	Interactive step-by-step reasoning; safe outputs	Susceptible to hallucination; reasoning not always reliable
GPT-4	2023	~1T (est., not public)	Multi-modal CoT, Tool-use, Self-consistency	Strong multi-step reasoning; fewer hallucinations	Closed weights; expensive inference
Claude (Anthropic)	2023	~50–100B (est.)	Constitutional AI + CoT	Ethical reasoning; interpretable rationales	Conservative answers; limited in math-heavy reasoning
Bard (Google)	2023	PaLM-based	CoT, real-time search augmentation	Access to real-time information; good general reasoning	Less accurate in formal logic or deep technical reasoning
Kosmos-1	2023	Not disclosed	Multi-modal CoT with visual grounding	Integrates vision-language reasoning	Still developing; unclear performance on pure logic tasks

CoT = Chain-of-Thought Prompting. **Tool Use** = Integration with external modules like calculators or code interpreters. **Instruction Tuning** = Fine-tuning with natural instructions and feedback. **Multi-modal** = Processes text, images, or other inputs together

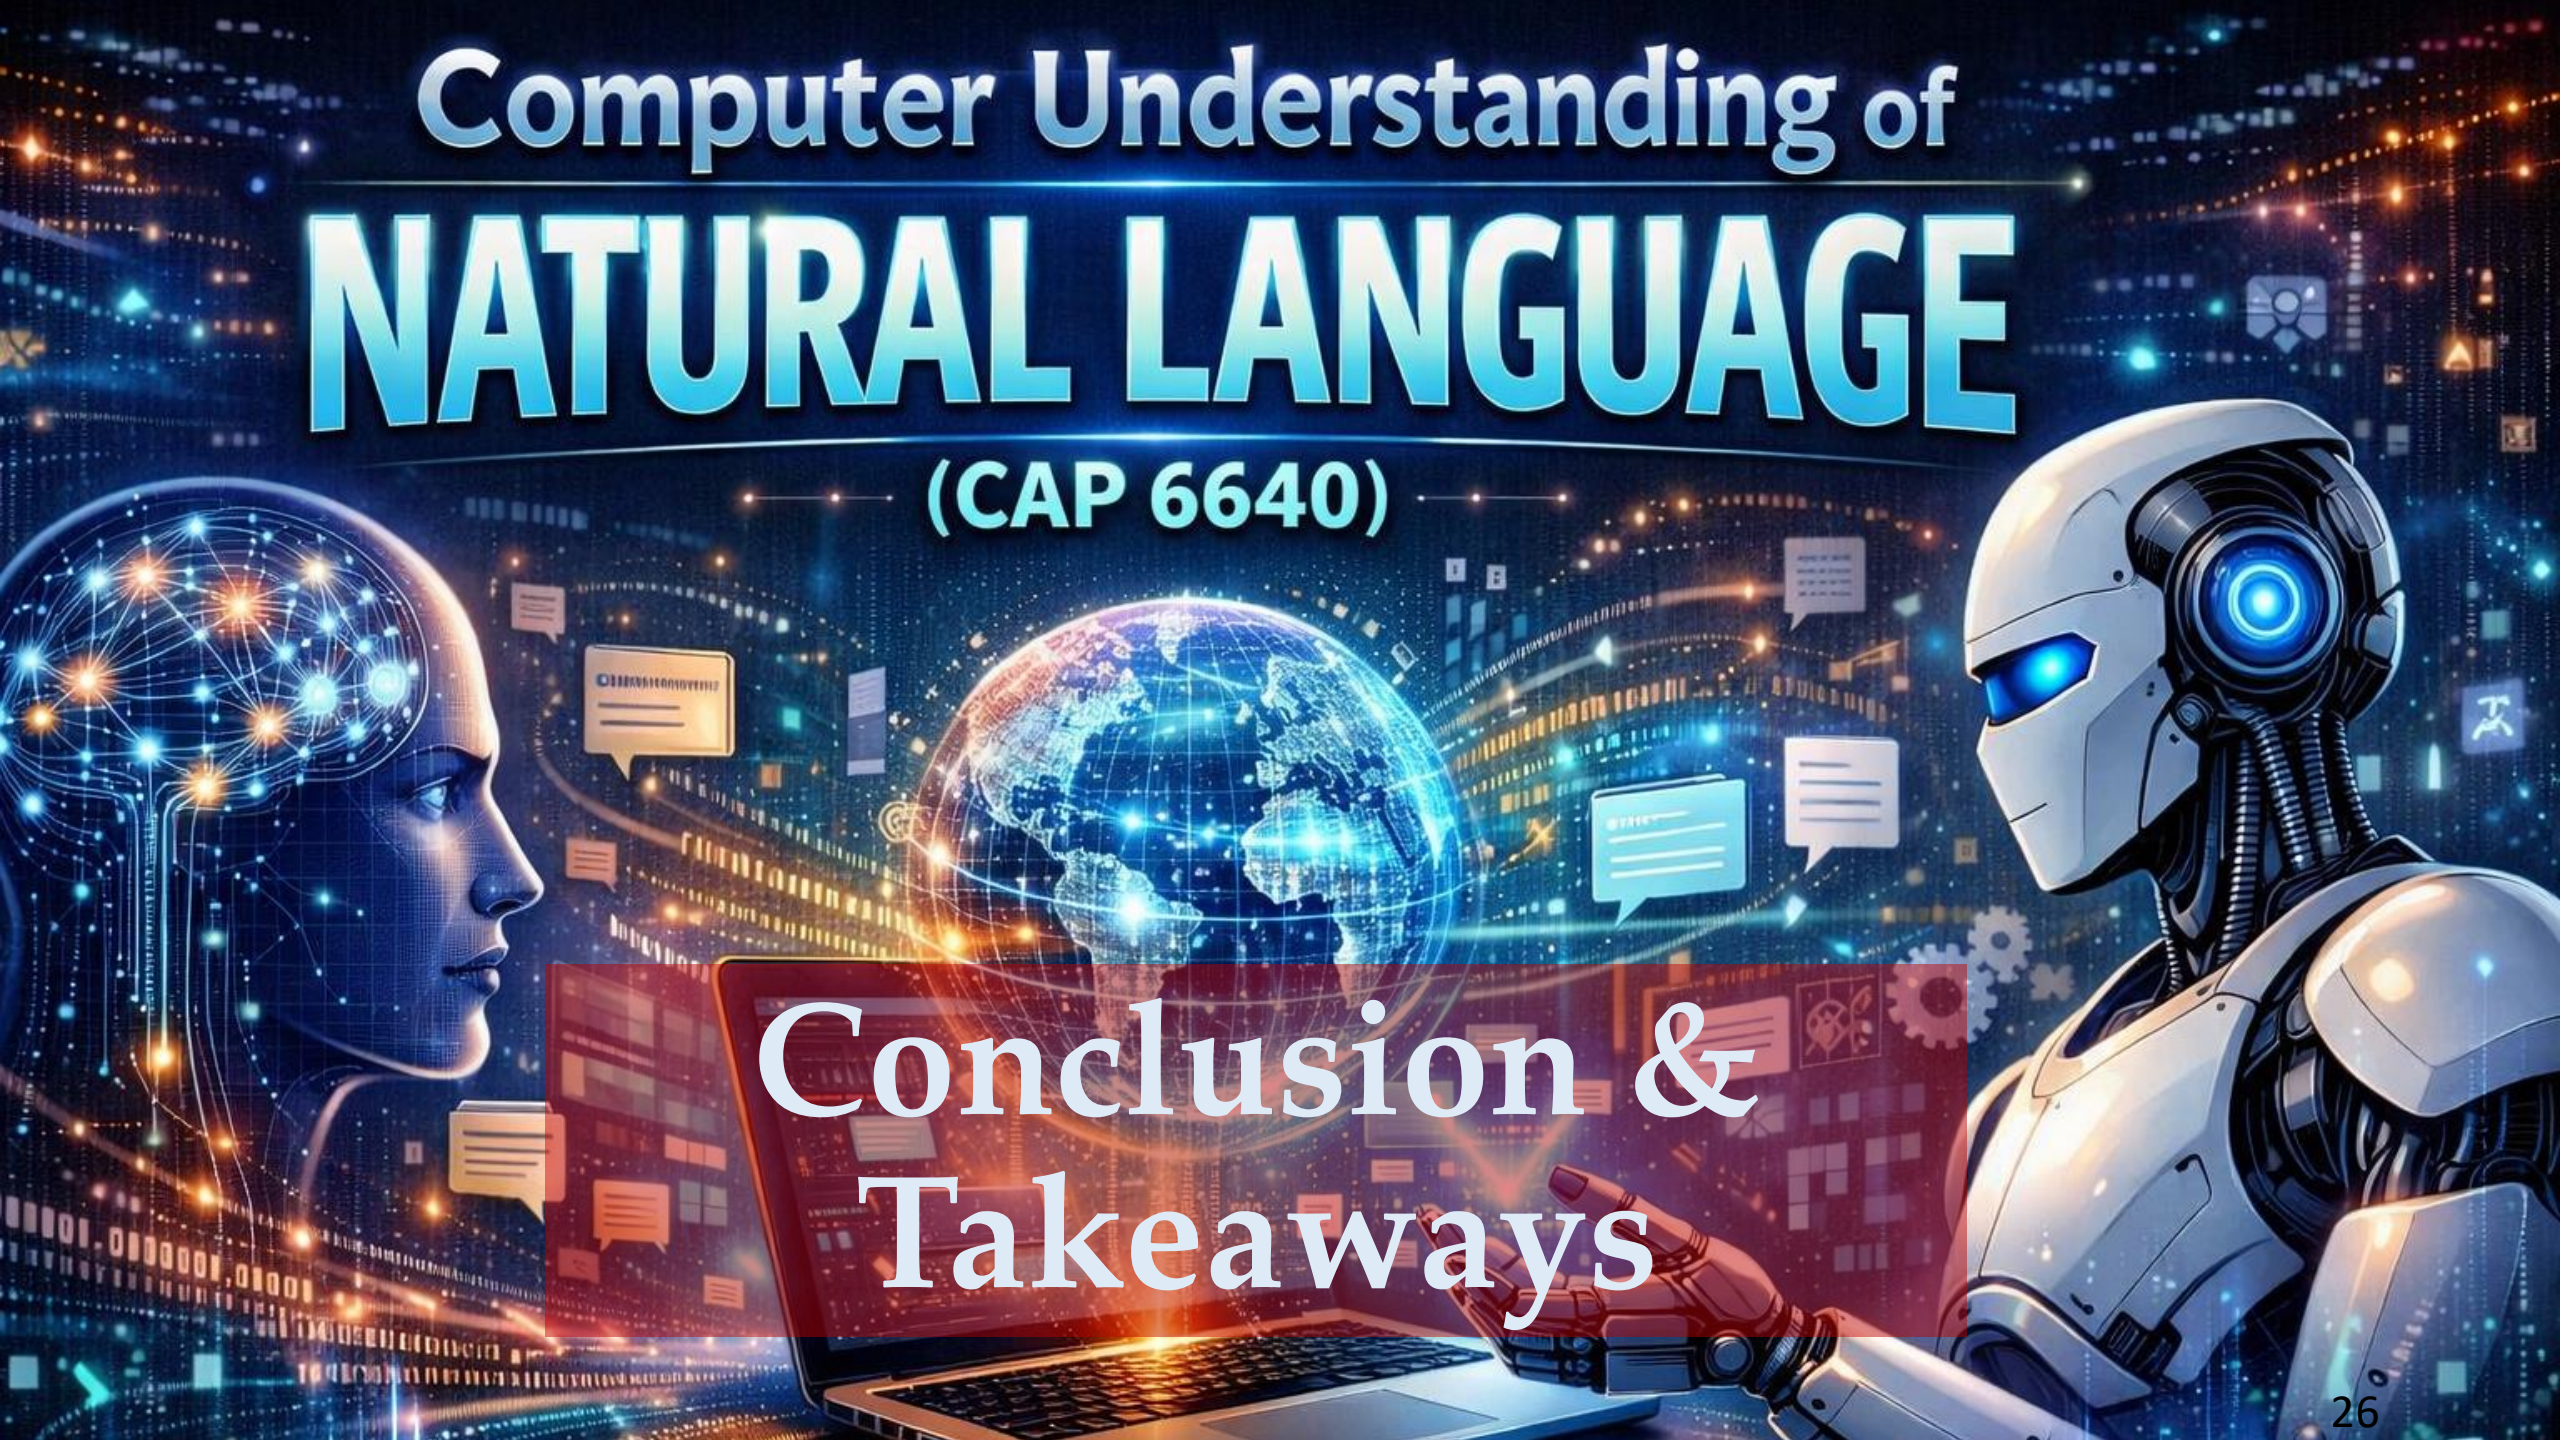

Conclusion & Takeaways

Conclusion & Takeaways

- **Humans** outperform machines in few-shot learning, generalization, and explainability.
- **LLMs** exhibit signs of informal reasoning, especially with techniques like **Chain-of-Thought (CoT)** prompting.
- High benchmark performance (e.g., PaLM, GPT-4) suggests **emergent reasoning**, but lacks robustness and symbolic grounding.
- Reasoning remains **pattern-based**, not truly logical or causal.
- Ongoing challenges: **definition, reliability, and evaluation** of reasoning in AI.
- Future directions: **neuro-symbolic models**, **tool-augmented reasoning**, and **explainable, verifiable outputs**.